

Computation-Performance Optimization of Convolutional Neural Networks with Redundant Kernel Removal

Chih-Ting Liu, Yi-Heng Wu, Yu-Sheng Lin, and Shao-Yi Chien

Media IC and System Lab

Graduate Institute of Electronics Engineering


National Taiwan University, Taiwan



Outline

- Introduction
 - Challenge of IOT+AI
- Pruning
- Proposed Method
 - Layer-wise Kernel Removal
 - Computation-Performance Optimization
- Experimental Results
- Conclusion

Internet of Things (IoT)




Internet of Things (IoT) + Artificial Intelligence (AI)

Now:
Intelligence on Cloud and Server
(GPU clusters)




Future:
Intelligence on Edge Devices
(CPU, ASIC)




IOT + Edge AI

- For visual tasks, Convolutional Neural Networks (**CNN**) is one of the AI solutions.
- Apply the CNN inference procedure on devices.
- Challenges
 - # **Parameters** of deep CNN.
(138M for VGG-16)
 - # **Operation** of deep CNN.
(Convolution layers)
→ Hard to implement on CPU or ASIC
Cost huge (1) **data access time** and
(2) **computation time**



Goal

- Can we analyze and remove the possible redundancy of Deep CNN models?



Picture: <https://goo.gl/images/oqYcfb>

6

Pruning– Related Works

- Parameters-pruning based method [1]
 - Remove the parameters with small **magnitude**.
 - Cause sparse weight matrix with same size.

→ Reduce parameters. But can **Not** reduce operations!!
- Kernel-pruning based method [2]
 - Remove the entire kernel (filter) when $\text{sparsity} > \text{threshold}$.


→ **Can** reduce operations! (reduction of convolution kernel)

[1] Han, Song, et al. "Learning both weights and connections for efficient neural network." Advances in neural information processing systems. 2015.

[2] C.-F. Chen, G. G. Lee, V. Sritapan, and C.-Y. Lin, "Deep convolutional neural network on iOS mobile devices," in Proc. 2016 IEEE International Workshop on Signal Processing Systems (SiPS).


Proposed Method

- Based on [2] , we make up the disadvantage of manually determining the threshold when removing kernels.
- Propose (1) Layer-wise Kernel Removal (LKR)
(2) Computation-Performance Optimization Flow (CPO)




Layer-wise Kernel Removal

For a specific convolutional layer l ,




Layer-wise Kernel Removal

For a specific convolutional layer l ,




Layer-wise Kernel Removal

For a specific convolutional layer l ,



Layer-wise Kernel Removal

For a specific convolutional layer l ,




Define **reducing factor** r_l , $0 \leq r_l \leq 1$:
The proportion of removed kernel.

Given r_l , we will remove Nr_l kernels.


Layer-wise Kernel Removal

For a specific convolutional layer l ,




Layer-wise Kernel Removal

For a specific convolutional layer l ,



Layer-wise Kernel Removal

For a specific convolutional layer l ,




Summary of LKR

- Given a **reducing factor** r_l for each layer, we can prune the Conv. layers. ($N - Nr_l$ kernels left)
- **Fast and Easy.**
- If we train from scratch with **smaller** network, not use pruning method, it may take time and may not train successfully.


Computation-Performance Optimization (CPO) Flow

- Monitoring the Performance to adjust the Computation Reduction.




CPO Flow

- Step 1 – Uniform Removal (UR)
 - Set **equal** r_l for every layer, obtain the performance drop.
 - Iteratively take experiments on **different** uniform r_l .
 - Pick an **acceptable** r_l^* among the experiments above.
 - * “**acceptable**” depends on the performance requirement by user.
 - Go to step2 to further adjust r_l^* for every layer.




CPO Flow

- Step 2 – Probing Model Sensitivity
 - Split model into few parts (ex. front , middle ,end)
 - Under same parameters left, find which part is less sensitive → **Remove more** on the part with **lower sensitivity**.
 - Increase r_l^* for one part to probe sensitivity.
 - decrease r_l^* for other parts to maintain params left.
 - Observe the performance drop based on diff. probing.
 - Choose the new reducing factor with smaller performance drop!



Experimental Results

- CNN model: **VDSR** (Very Deep Super Resolution) model [3]
 - 20 layer fully convolutional Network → Excluding the influence of FC layers.



ILR: Interpolated Low Resolution Image

HR: High Resolution Image

- Retrain Epoch: 8
- Evaluation Dataset: (1) Set5 X2 (2) Set14 X2
- Performance Evaluation : **PSNR (dB)**

[3] Kim et al. "Accurate image super-resolution using very deep convolutional networks." CVPR 2016.

CPO step 1 Result – Uniform Removal(UR)


Model	Kernel per layer	Set5 PSNR	Set14 PSNR	Parameters	FLOP
Original	64	37.50 dB	33.08 dB	6.7×10^5	1.11×10^9

Reducing Factor (r)	Kernel per layer	Set5 PSNR Drop ↓ (dB)	Set14 PSNR Drop ↓ (dB)	Params Remained	FLOP Remained
0.00	64	0	0	100 (%)	100 (%)
0.12	56	0.13	0.19	76.60 (%)	76.58(%)
0.19	52	0.21	0.21	66.07 (%)	66.04(%)
0.25	48	0.19	0.25	56.32 (%)	56.28(%)
0.31	44	0.20	0.25	47.34 (%)	47.30(%)
0.38	40	0.36	0.36	39.15 (%)	39.10(%)
0.44	36	0.40	0.39	31.73 (%)	31.68(%)



CPO step 2 – Probe Model Sensitivity

- Split the model into three parts: **Front, Middle ,End**



- Start from $r^* = 0.25$, and adjust the value for different parts.
- Try to **improve the performance** by pruning the **less sensitive part more** under almost same computation constraint.

CPO step2 Results

Reducing Factor (r) (F , M , E)	Kernel per parts (F, M, E)	Set5 PSNR Drop ↓ (dB)	Set14 PSNR Drop ↓ (dB)	Parameters Remained (%)
(0.25,0.25,0.25)	48, 48, 48	0.19	0.25	56.32 (UR)
(0.44,0.18,0.18)	36, 52, 52	0.19	0.23	55.39
(0.44,0.12,0.25)	36, 56, 48	0.18	0.20	56.36
(0.12,0.18,0.44)	56, 52, 36	0.24	0.29	58.60
(0.18,0.18,0.38)	52, 52, 40	0.30	0.26	57.54
(0.25,0.44,0.06)	48, 36, 60	0.27	0.30	55.94
(0.18,0.44,0.16)	52, 36, 56	0.26	0.30	55.51

→ The **Front part** of VDSR model is less sensitive; therefore, we can **increase r_{front}** .

→ Use (0.44,0.12,0.25) to prune VDSR can achieve **50%** parameters removal and improve the UR performance. (PSNR drop 0.18/0.20)

Conclusion

- Layer-wise kernel removal can remove the redundancy of CNN and at the meantime reduce operations.
- Using Computation-Performance Optimization (CPO) Flow can remove more kernels in the model parts with more redundancy than others.

Extension

- CPO now is empirical, we've extended this work to develop a more **systematic CPO flow**.
- Have performed new CPO on VGG-19 for Cifar-10 classification.
- Compared with the state-of-the-art kernel removal method.
- Have submitted our extended work to **TCAS-1**.



Q & A

Thanks for Listening!

Backup

■ Pruning on VGG-19 of Cifar-10

	Reducing Factor r	Drop(%)	Params Remained	FLOP Remained
Original	0	0%	2.0×10^7	4.0×10^8
		(Acc: 92.19%)	100%	100%
Uniform	0.125	0.10%	77.86%	76.87%
Removal	0.250	0.54%	58.47%	56.78%
(UR) [13]	0.375	1.36%	41.83%	39.72%
	0.500	3.23%	27.96%	25.70%
	0.625	6.42%	16.83%	14.72%

	D_{exp}	Final Drop	Params Remained	FLOP Remained
CPO	0.10%	0.06%	26.89%	51.99%
(ours)	0.54%	0.50%	9.80%	35.02%
	1.36%	1.35%	5.16%	25.77%
	3.23%	3.02%	3.42%	20.96%
	6.42%	6.30%	1.76%	16.76%

Backup

- VGG-16 on Cifar-10
- Compare with state-of-the-art [4]

Model	Error	FLOP	Pruned	Params	Pruned
VGG-16 [11]	6.75%	3.13×10^8		1.5×10^7	
Pruned [11]	6.60%	2.06×10^8	34.2%	5.4×10^6	64.00%
VGG-16	6.22%	3.13×10^8		1.5×10^7	
Pruned(CPO)	6.16%	2.09×10^8	33.4%	3.2×10^6	78.76%

[4] H. Li, A. Kadav, I. Durdanovic, H. Samet, and H. P. Graf, “Pruning filters for efficient convnets,” arXiv preprint arXiv:1608.08710, 2016